

Meta-Learning

— an idiosyncratic tutorial

Yee Whye Teh

Statistics @ Oxford

DeepMind

<http://csml.stats.ox.ac.uk/people/teh/>



Yee Whye Teh


Small Data Problems

- Few Shot Learning




Small Data Problems

- Recommender Systems



Small Data Problems

- Robotics




Small Data Problems

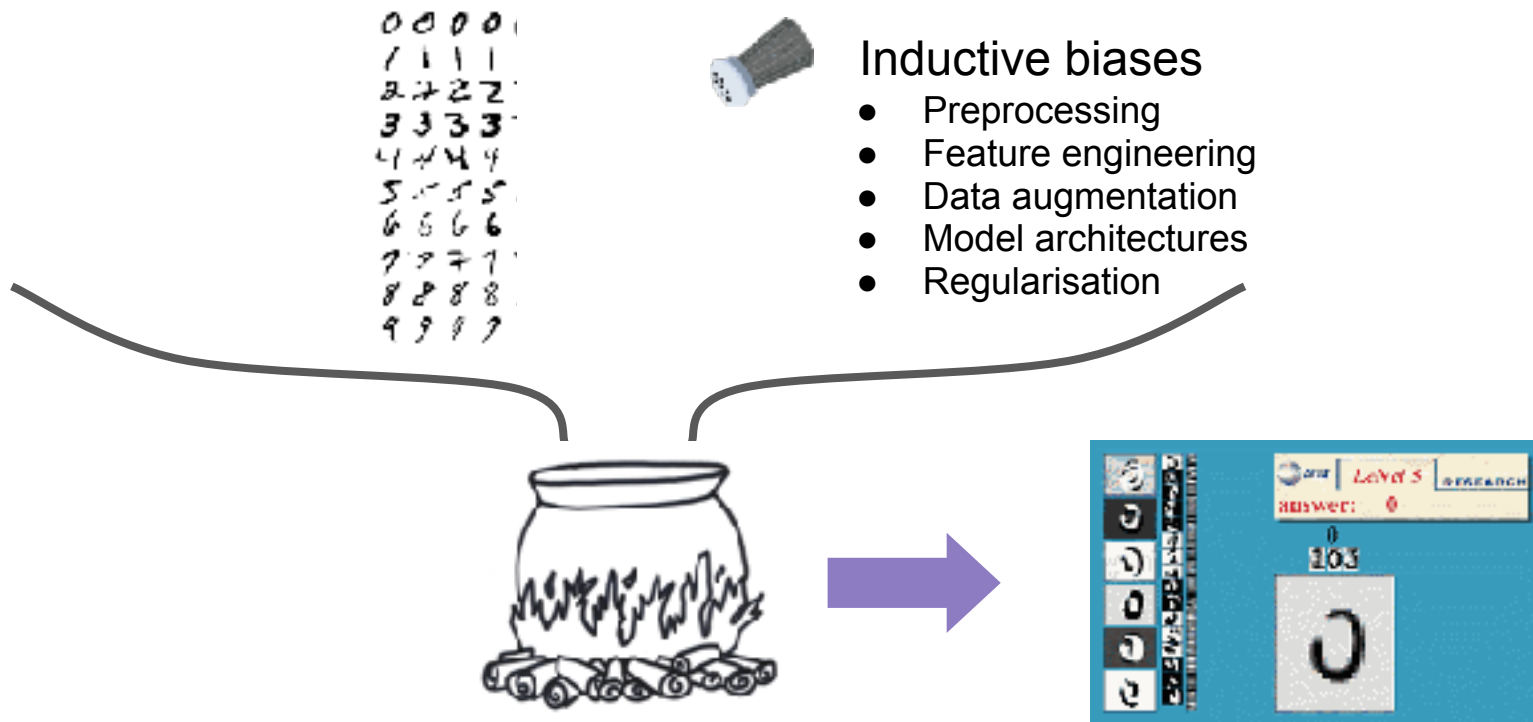
- Artificial General Intelligence

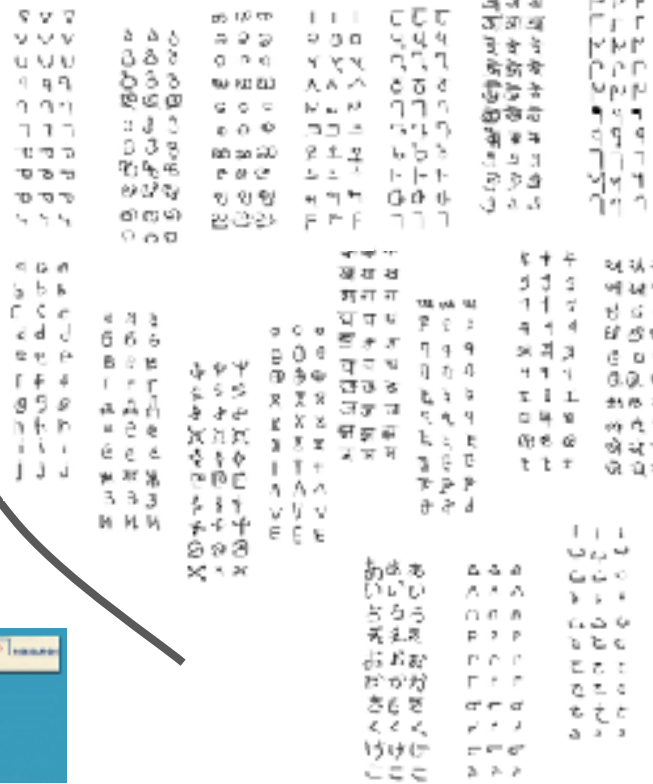


Machine Learning with Big Data



Machine Learning with Small Data



[illegible]

Meta-Learning: an idiosyncratic tutorial

- Optimisation perspective on meta-learning
 - Optimisation-based meta-learning
 - Black-box meta-learning
- Probabilistic perspective on meta-learning
 - Stochastic processes
 - Neural processes
 - Uncertainty in meta-learning
- Probabilistic symmetries and neural architectures
- Note: no meta reinforcement learning (meta-RL)

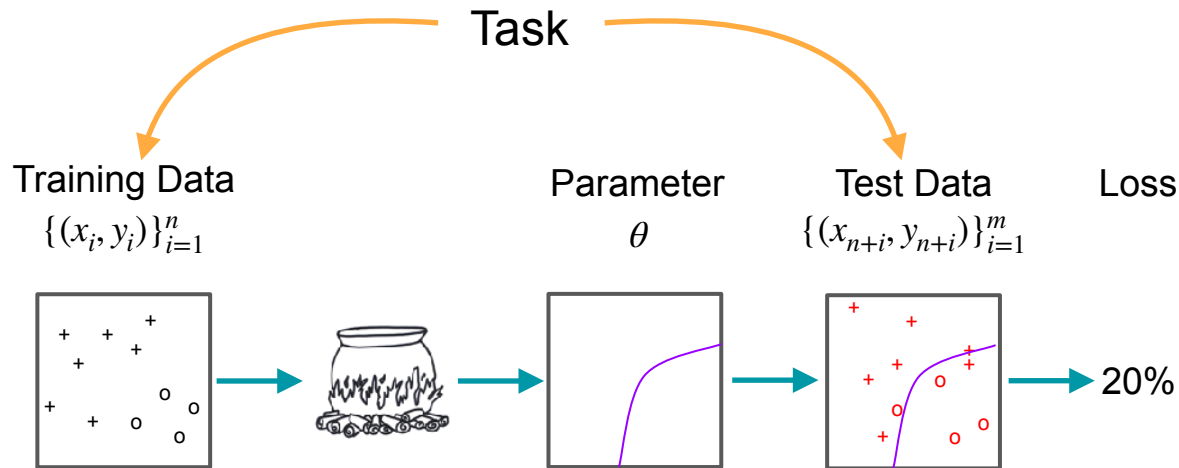


Meta-Learning: an idiosyncratic tutorial

- **Optimisation perspective on meta-learning**
 - Optimisation-based meta-learning
 - Black-box meta-learning
- Probabilistic perspective on meta-learning
 - Stochastic processes
 - Neural processes
 - Uncertainty in meta-learning
- Probabilistic symmetries and neural architectures
- Note: no meta reinforcement learning (meta-RL)



Single-Task Learning

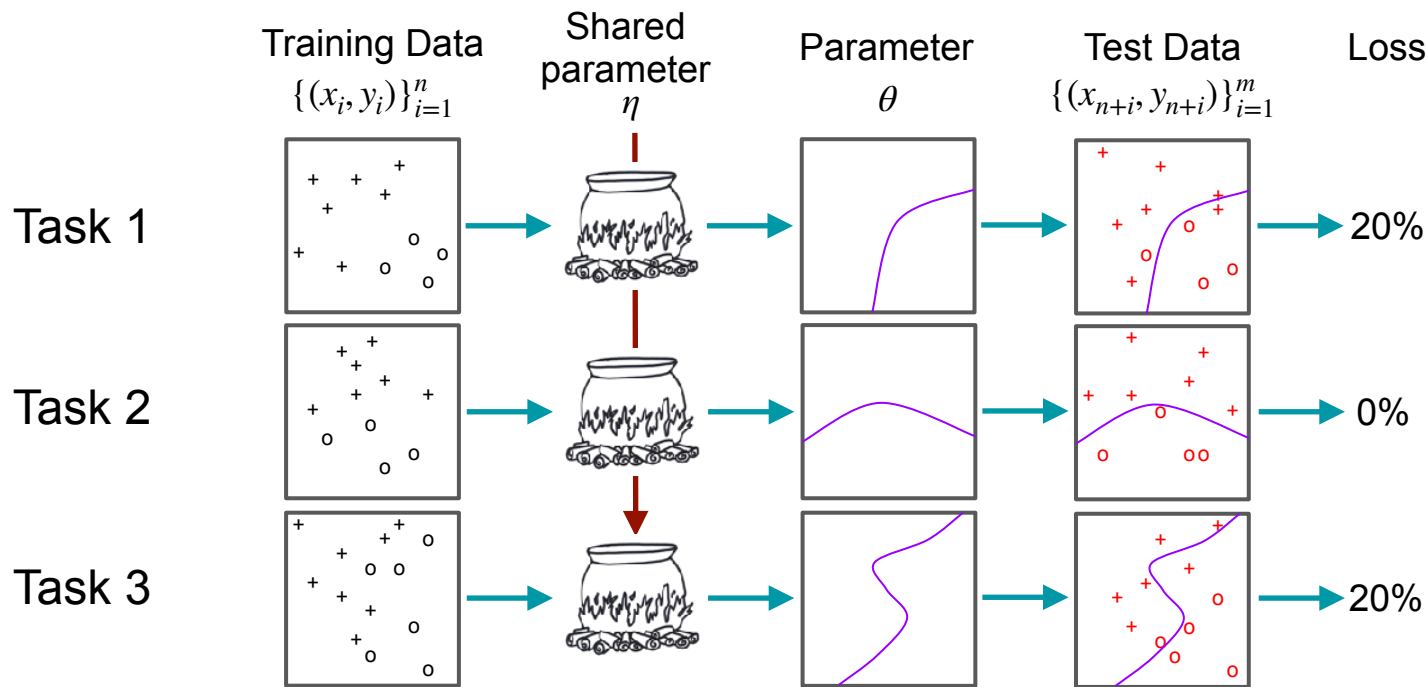


Empirical Risk Minimisation

$$\arg \min_{\theta} \sum_{i=1}^n \text{Loss}(f_{\eta, \theta}(x_i), y_i) + \text{Regulariser}(\theta)$$



Multi-Task Learning



$$\arg \min_{\eta, \{\theta_j\}} \sum_{j=1}^T \sum_{i=1}^n \text{Loss}(f_{\eta, \theta_j}(x_{ji}), y_{ji}) + \text{Regulariser}(\theta_j)$$



Meta-Learning and Learning-to-Learn

- Is it possible to learn to generalise from training to test data?

- Learn:

$$\theta_j = \arg \min_{\theta_j} \sum_{i=1}^n \text{Loss}(f_{\eta, \theta_j}(x_{ji}), y_{ji}) =: \text{Learner}(\eta, \text{TrainData}_j)$$

- Test performance:

$$\sum_{i=n+1}^{n+m} \text{Loss}(f_{\eta, \theta_j}(x_{ji}), y_{ji}) =: \mathcal{L}(\eta, \theta_j, \text{TestData}_j)$$

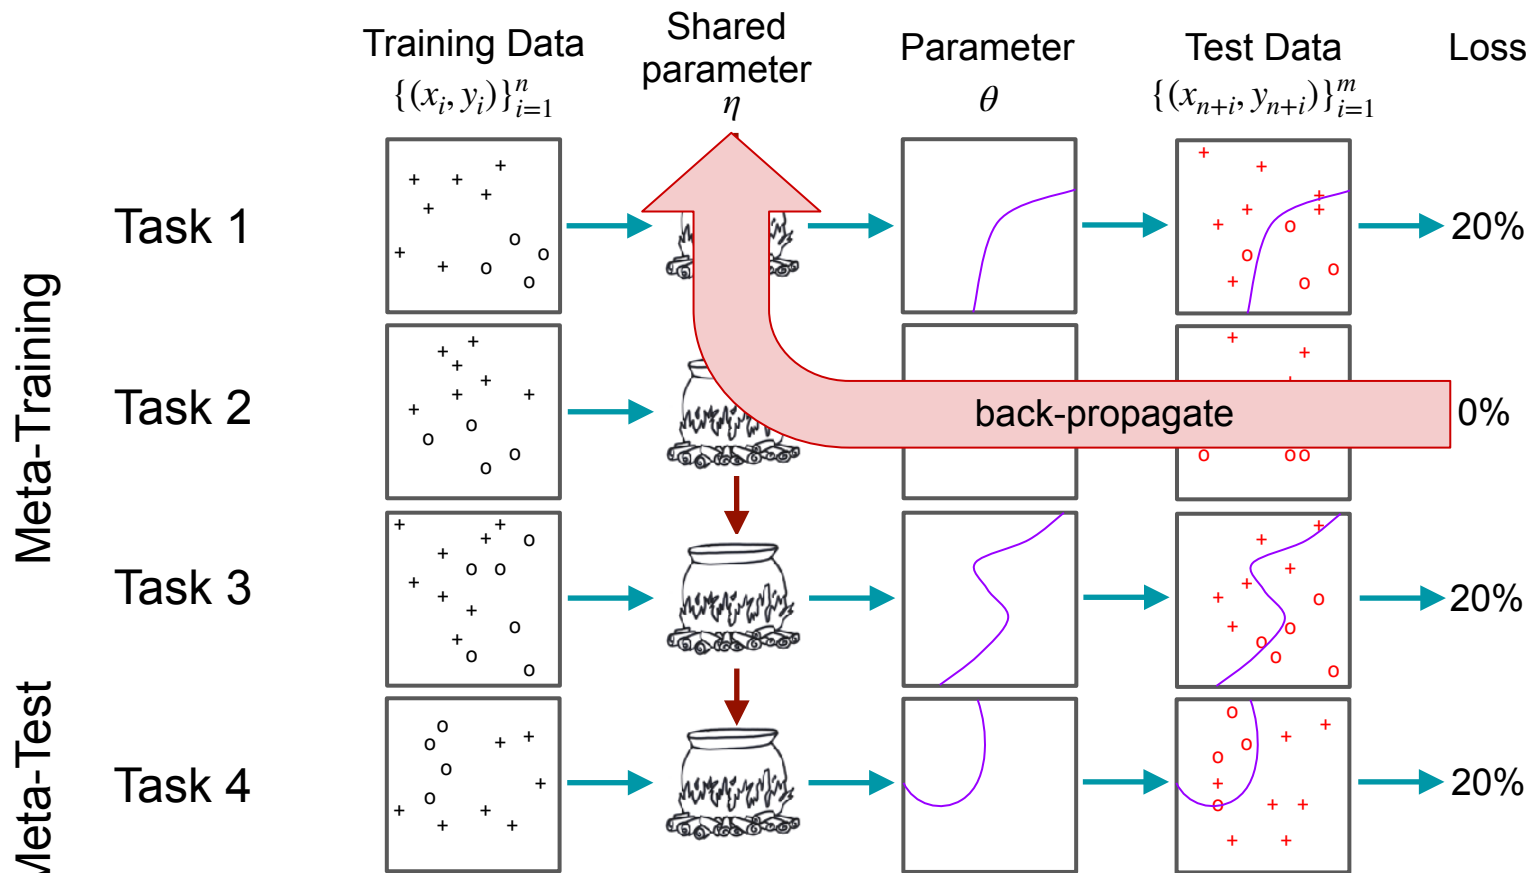
- Optimise shared parameters for test performance:

$$\arg \min_{\eta} \sum_{j=1}^T \mathcal{L}(\eta, \theta_j, \text{TestData}_j) = \arg \min_{\eta} \sum_{j=1}^T \sum_{i=n+1}^{n+m} \text{Loss}(f_{\eta, \text{Learner}(\eta, \text{TrainData}_j)}(x_{ji}), y_{ji})$$

- “Our training procedure is based on a simple machine learning principle: test and train conditions must match” — [[Vinyals et al NeurIPS 2016](#)]



Meta-Learning



Meta-Learning

- For each iteration:

- Pick (minibatch of) training task j :

- Base learner:

$$\theta_j = \text{Learner}(\eta, \text{TrainData}_j) = \arg \min_{\theta_j} \sum_{i=1}^n L(f_{\eta, \theta_j}(x_{ji}), y_{ji})$$

- Test performance:

$$\mathcal{L}(\eta, \theta_j, \text{TestData}_j) = \sum_{i=n+1}^{n+m} \text{Loss}(f_{\eta, \theta_j}(x_{ji}), y_{ji})$$


- Meta-learner: optimise shared parameters for test performance:

$$\eta \leftarrow \eta - \epsilon \frac{d}{d\eta} \sum_{i=n+1}^{n+m} \text{Loss}(f_{\eta, \text{Learner}(\eta, \text{TrainData}_j)}(x_{ji}), y_{ji})$$




Image Datasets

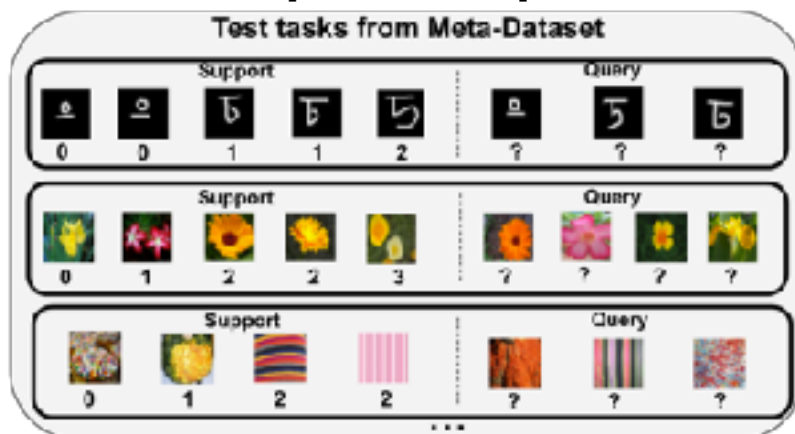
[[Omniglot](#)]



[[mini-ImageNet](#)]



[[Meta Dataset](#)]



Dataset (other names)

ilsvrc_2012 (ImageNet, ILSVRC) [[instructions](#)]

omniglot [[instructions](#)]

aircraft (FGVC-Aircraft) [[instructions](#)]

cu_birds (Birds, CUB-200-2011) [[instructions](#)]

dtd (Describable Textures, DTD) [[instructions](#)]

quickdraw (Quick, Draw!) [[instructions](#)]

fungi (FGVCx Fungi) [[instructions](#)]

vgg_flower (VGG Flower) [[instructions](#)]

traffic_sign (Traffic Signs, German Traffic Sign Recognition Benchmark, GTSRB) [[instructions](#)]

mscoco (Common Objects in Context, COCO) [[instructions](#)]

Total (All datasets)




Different Kinds of Meta-Learning Methods

- Different choices of **Learner** and function class f lead to different methods.
 - Optimisation-based
 - black-box
 - metric-based
 - memory-based
 - Bayesian...
- Different choices correspond to different **inductive biases**.
 - Psychological theories
 - Symmetries, invariances and equivariances
 - Computational and learnability constraints
 - Intuitions and experimental findings



Optimisation-based Meta-Learning



- Optimisation-based base learner.

θ_0

$$\theta_1 \leftarrow \theta_0 + \text{Update}_{\eta}(\nabla_{\theta_0} L(f_{\theta_0}(X_1), Y_1))$$

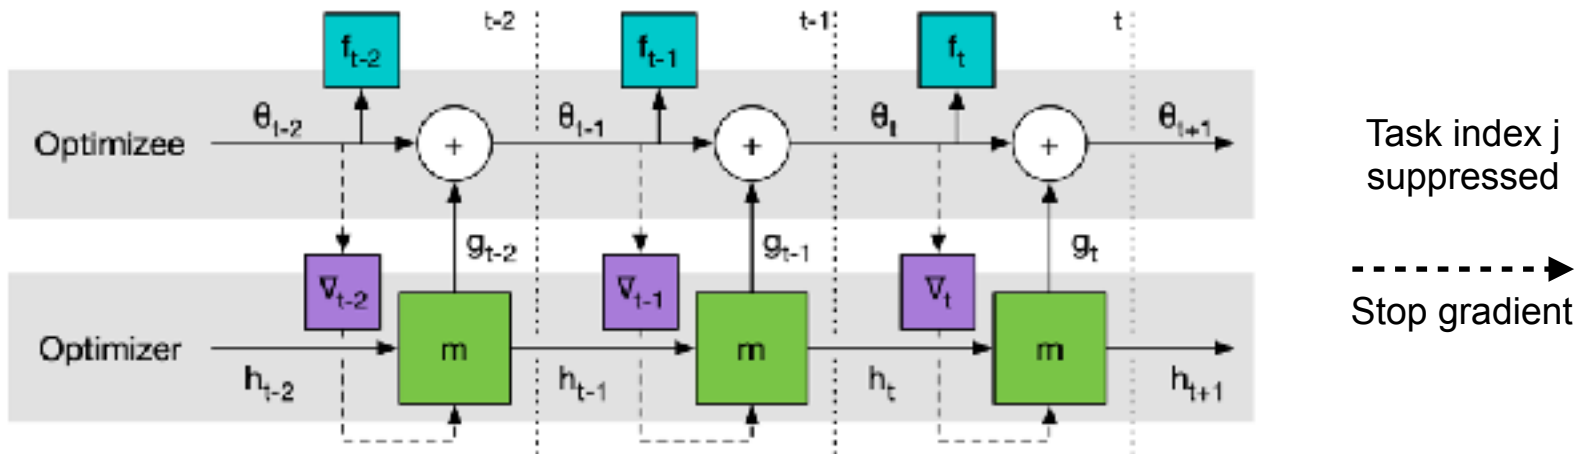
$$\theta_2 \leftarrow \theta_1 + \text{Update}_{\eta}(\nabla_{\theta_1} L(f_{\theta_1}(X_2), Y_2))$$

\vdots

- Meta-learn meta-parameters θ_0, η .



Learning to Learn by Gradient Descent by Gradient Descent & LSTM Meta-Learner



$$f_t = f_{\theta_t}(X_t)$$

$$\nabla_t = \frac{d}{d\theta_t} L(f_t, Y_t)$$

X_t : training inputs in iter t
 Y_t : training targets in iter t

ADAM

$$\mu_t = \beta_1 \mu_{t-1} + (1 - \beta_1) \nabla_t$$

$$V_t = \beta_2 V_{t-1} + (1 - \beta_2) \nabla_t^2$$

$$g_t = - \frac{\epsilon}{\sqrt{V_t} + \delta} \mu_t$$

$$\theta_{t+1} = \theta_t + g_t$$



Model-Agnostic Meta-Learning (MAML)


Single (or few) step gradient descent learner:

$$\theta_j = \eta - \epsilon \nabla_{\eta} L(f_{\eta}(X_j), Y_j)$$

Meta-learn shared initialisation η :

$$\eta \leftarrow \eta - \epsilon_0 \frac{d}{d\eta} \sum_j L(f_{\theta_j}(X_{jt}), Y_{jt})$$

Compute gradient through learner gradient step.



Model-Agnostic Meta-Learning (MAML)

$$\begin{aligned}\theta_j &= \eta - \epsilon \nabla_{\eta} L(f_{\eta}(X_j), Y_j) \\ \frac{d}{d\eta} L(f_{\theta_j}(X_{jt}), Y_{jt}) &= \nabla_{\theta_j} L(f_{\theta_j}(X_{jt}), Y_{jt}) \frac{d\theta_j}{d\eta} \\ &= \nabla_{\theta_j} L(f_{\theta_j}(X_{jt}), Y_{jt}) \left(I - \epsilon \nabla_{\eta}^2 L(f_{\eta}(X_{jt}), Y_{jt}) \right)\end{aligned}$$

- Difficulty is in computing $\nabla_{\theta_j} L(f_{\theta_j}(X_{jt}), Y_{jt}) \nabla_{\eta}^2 L(f_{\eta}(X_{jt}), Y_{jt})$.
- Vector-Hessian products can be automatically computed in linear time.




Optimisation-based Meta-Learning

- Two level optimisation problem
 - Flexible and applicable to wide range of neural architectures
 - Positive inductive bias
- Challenges
 - sensitive to neural architectures
 - can be expensive and unstable
- Difficulties of back-propagating through many base learner iterations.
 - If base learner optimisation can be analytically solved, gradients can be computed exactly [[Harrison et al WAFC 2018](#), [Lee et al CVPR 2019](#)].
 - Alternatively, use implicit function theorem [[Rajeswaran et al NeurIPS 2019](#)].




Black-box Meta-Learning

- So far, the base learner is an **optimisation-based** learner.
- Not all learning algorithms are optimisation-based.
- For example, in k -nearest neighbour, learner simply “memorises” training data, and matches test inputs with memory to make predictions.
- **Black-box meta-learning:** Implement the base learner using a differentiable programming framework.



Matching Networks




- Embed each train input $x_i \mapsto g_\eta(x_i)$.
- Embed test input $x_i \mapsto f_\eta(x_i)$

- “Softened” 1-NN classifier:

$$p(\cdot | x_{n+j}, \text{TrainData}) = \sum_{i=1}^n \frac{e^{g_\eta(x_i)^\top f_\eta(x_{n+j})}}{\sum_{l=1}^n e^{g_\eta(x_l)^\top f_\eta(x_{n+j})}} y_i$$

- Memory-based meta-learning
- Metric-based meta-learning

Prototypical Networks



- Embed each input $x_i \mapsto f_\eta(x_i)$.

- Form “prototypes”:


$$c_k = \frac{\sum_{i:y_i=k} f_\eta(x_i)}{\sum_{i:y_i=k} 1}$$

- Predict:

$$p(y_{n_j} = k | x_{n+j}, \text{TrainData}) = \frac{e^{-\|f_\eta(x_{n+j}) - c_k\|^2 / \sigma^2}}{\sum_{l=1}^K e^{-\|f_\eta(x_{n+j}) - c_l\|^2 / \sigma^2}}$$



Memory-Augmented Neural Networks (MANNs)



Each task: sequence of input/output pairs $(x_1, y_1), (x_2, y_2), \dots$


Base learner directly predicts

$$p_{\eta}(y_t | (x_i, y_i)_{i=1}^{t-1}, x_t)$$

Neural Turing machine with external memory.



Simple Neural Attentive Learner (SNAIL)




$$p_{\eta}(y_t | (x_i, y_i)_{i=1}^{t-1}, x_t)$$

Simple:

- treats base learner problem as sequential prediction
- Use (causal) convolution and attention layers.



Base Learner Architecture and Task Representation

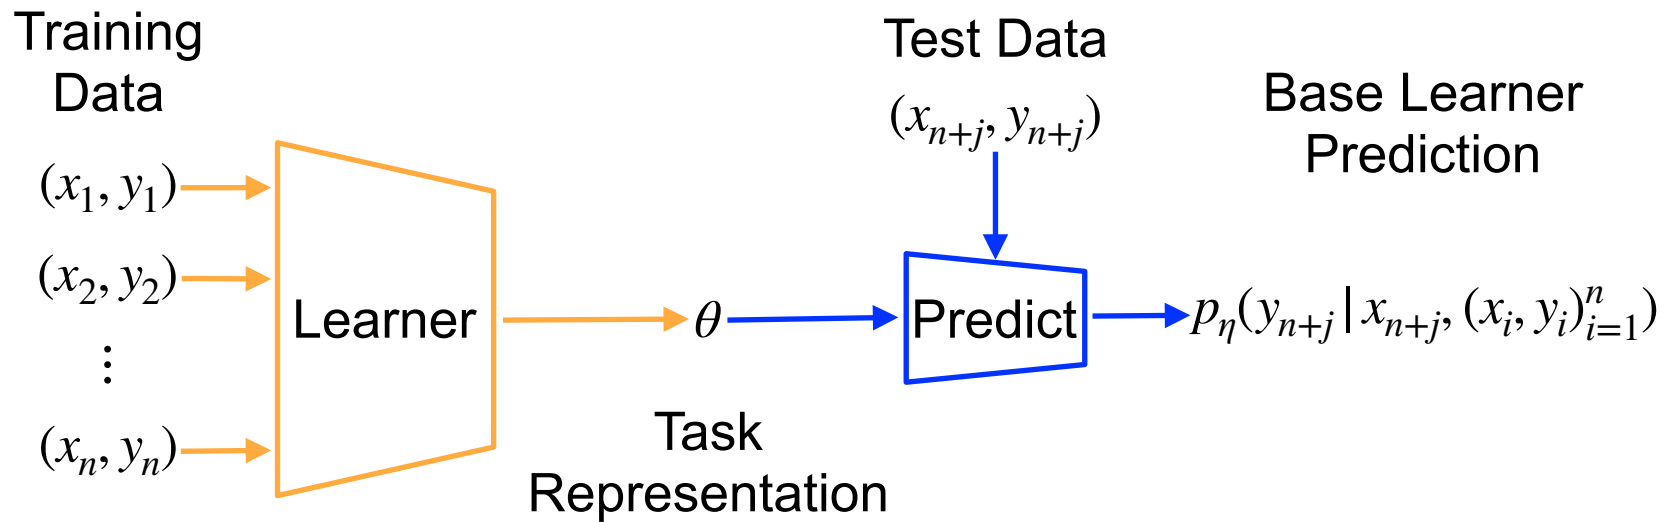


Black-box Meta-Learning

- Learn a differentiable function to map from training data to test predictions.
 - Simple: reduces meta-learning back to supervised learning.
 - Broad and flexible framework.
- Challenges:
 - Harder to learn as no inductive bias for base learner to optimise on training data.
 - Less able to generalise out of meta-training distribution.



Base Learner Architecture and Task Representation



Permutation Invariance

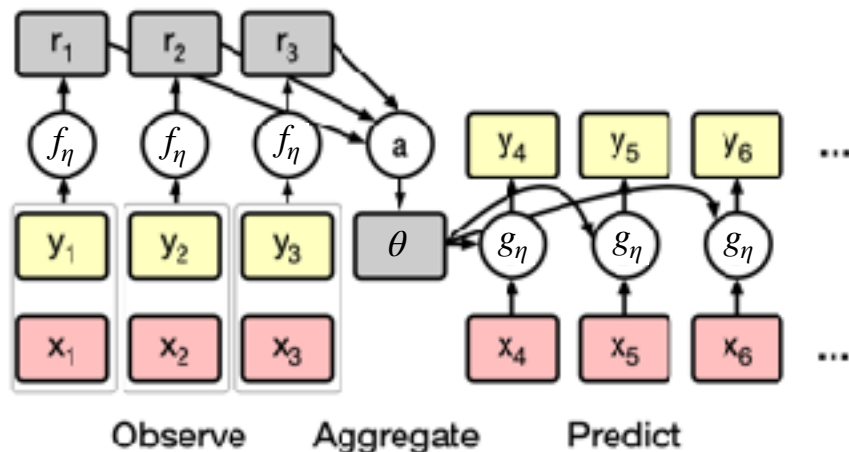
- Assumed that data items within each task are iid.
- Learner function should be invariant to permutations of the training data:

$$\text{Learner}(\eta, \{(x_1, y_1), \dots, (x_n, y_n)\}) = \text{Learner}(\eta, \{(x_{\pi(1)}, y_{\pi(1)}), \dots, (x_{\pi(n)}, y_{\pi(n)})\})$$

- MAML, Prototypical Nets (and simpler version of Matching Nets) are.
- LSTM meta-learner, MANN, SNAIL (and a version of Matching Nets) are not permutation invariant.
- Permutation invariance is an inductive bias.
- How to design neural architectures for permutation invariance? [more later]



Conditional Neural Processes



- Embed input/output pairs

$$(x_i, y_i) \mapsto f_\eta(x_i, y_i)$$

- Aggregate embeddings

$$\theta = \frac{1}{n} \sum_{i=1}^n f_\eta(x_i, y_i)$$

- Permutation invariant.
- Interpret θ as a task representation.

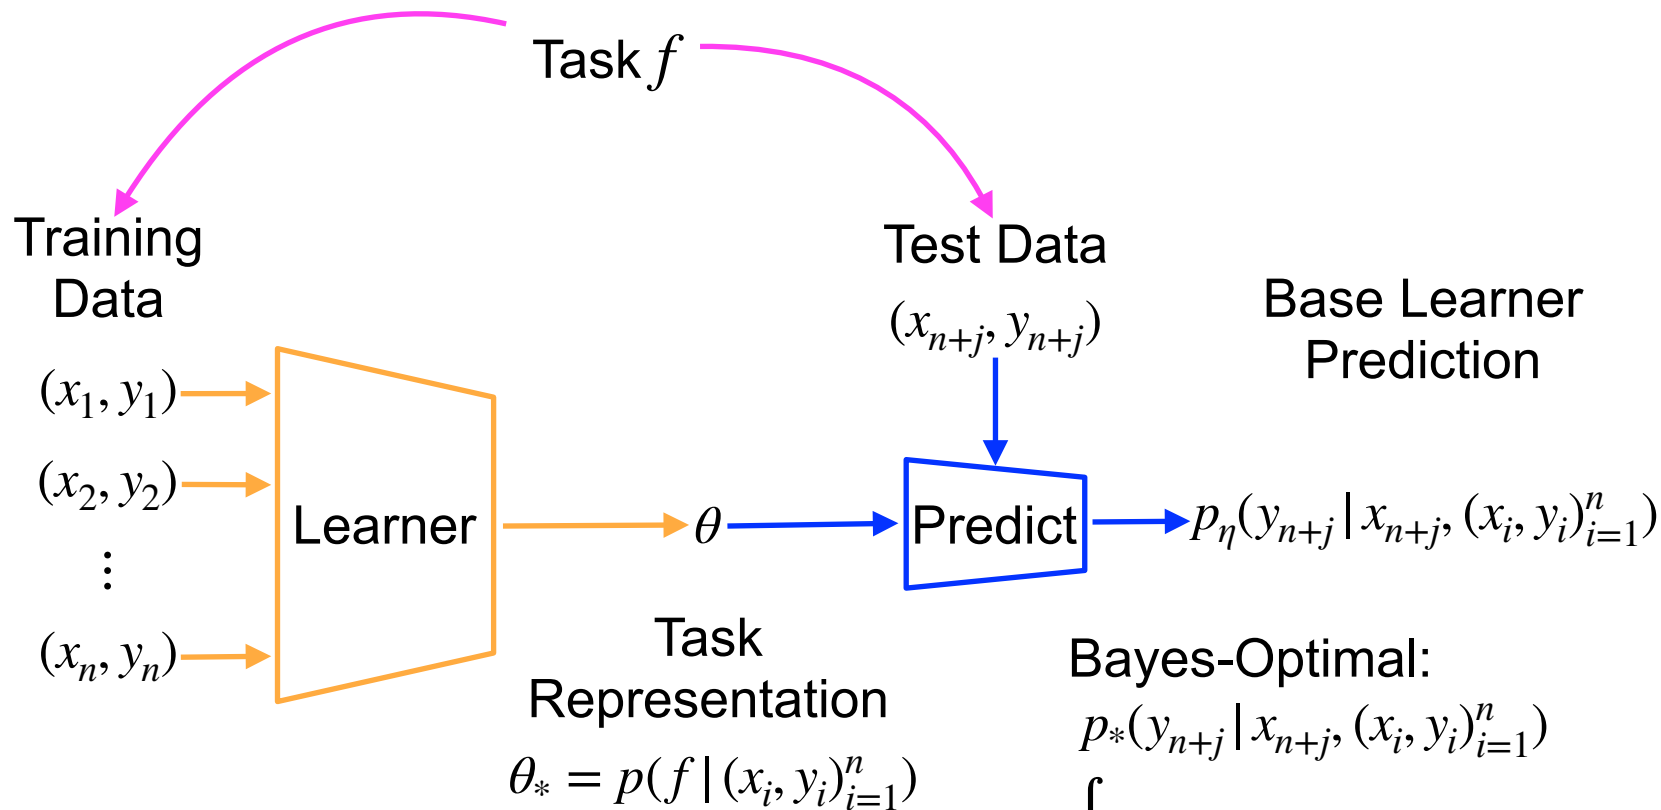


Meta-Learning: an idiosyncratic tutorial

- Optimisation perspective on meta-learning
 - Optimisation-based meta-learning
 - Black-box meta-learning
- **Probabilistic perspective on meta-learning**
 - Stochastic processes
 - Neural processes
 - Uncertainty in meta-learning
- Probabilistic symmetries and neural architectures
- Note: no meta reinforcement learning (meta-RL)



Probabilistic Perspective on Meta-Learning

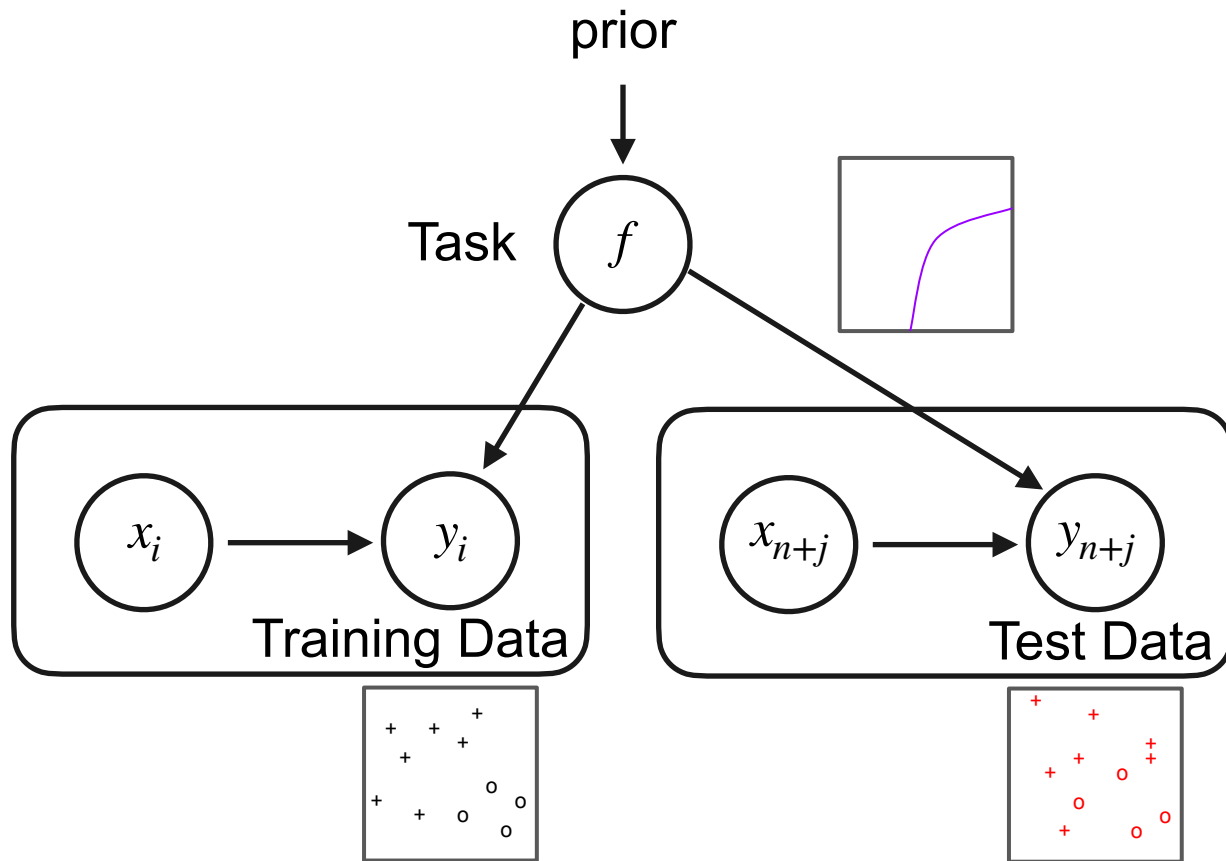


Bayes-Optimal:

$$p_*(y_{n+j} | x_{n+j}, (x_i, y_i)_{i=1}^n) \\ = \int p(y_{n+j} | x_{n+j}, f) p(f | (x_i, y_i)_{i=1}^n) df$$






Generative Process



Specifying Stochastic Processes

- Gaussian processes are typically described via marginal distributions:

$$\begin{pmatrix} f(x_1) \\ f(x_2) \\ \vdots \\ f(x_t) \end{pmatrix} \sim \mathcal{N} \left(\begin{pmatrix} \mu(x_1) \\ \mu(x_2) \\ \vdots \\ \mu(x_t) \end{pmatrix}, \begin{pmatrix} K(x_1, x_1) & K(x_1, x_2) & \cdots & K(x_1, x_t) \\ K(x_2, x_1) & K(x_2, x_2) & \cdots & K(x_2, x_t) \\ \vdots & \vdots & \ddots & \vdots \\ K(x_t, x_1) & K(x_t, x_2) & \cdots & K(x_t, x_t) \end{pmatrix} \right)$$



Specifying a Stochastic Process

- A stochastic process is a joint distribution over an infinite collection of random variables $(f(x))_{x \in \mathcal{X}}$.
- Kolmogorov Extension Theorem:
 - Constructs a stochastic process by specifying its finite dimensional marginal distributions.
 - Family of finite dimensional joint distributions $\rho_{x_{1:n}}$, one for each $n \in \mathbb{N}$ and finite sequence $x_{1:n} \in \mathcal{X}$. We will want these to form the marginals:

$$\rho_{x_{1:n}}(y_{1:n}) = p(f(x_1) = y_1, \dots, f(x_n) = y_n)$$



Exchangeability and Consistency

- Exchangeability: for each n , $x_{1:n}$, and permutation π of $\{1, \dots, n\}$

$$\rho_{x_1, \dots, x_n}(y_1, \dots, y_n) = \rho_{x_{\pi(1)}, \dots, x_{\pi(n)}}(y_{\pi(1)}, \dots, y_{\pi(n)})$$

- Consistency: for each n, m , $x_{1:n+m}$

$$\rho_{x_1, \dots, x_n}(y_1, \dots, y_n) = \int \rho_{x_{1:n+m}}(y_{1:n+m}) dy_{n+1:n+m}$$



Bayesian nonparametrics




- Gaussian processes: regression, classification [[Neal 1994](#), [Rasmussen & Williams 2006](#)]
- Dirichlet processes: infinite mixture models, clustering [[Neal JCGS 1999](#), [Rasmussen NeurIPS 2000](#)]
- Hierarchical Dirichlet processes: topic models and HMMs [[Teh et al JASA 2006](#)]
- Pitman-Yor and Poisson-Kingman processes: power laws and language models [[Teh ACL 2006](#)], species discovery problems [[Favaro et al Biometrics 2015](#)]
- Sparse networks models and power-law structures [[Caron & Fox JRSSB 2017](#)]



Specifying Stochastic Processes

- Gaussian processes are typically described via marginal distributions:

$$\begin{pmatrix} f(x_1) \\ f(x_2) \\ \vdots \\ f(x_t) \end{pmatrix} \sim \mathcal{N} \left(\begin{pmatrix} \mu(x_1) \\ \mu(x_2) \\ \vdots \\ \mu(x_t) \end{pmatrix}, \begin{pmatrix} K(x_1, x_1) & K(x_1, x_2) & \cdots & K(x_1, x_t) \\ K(x_2, x_1) & K(x_2, x_2) & \cdots & K(x_2, x_t) \\ \vdots & \vdots & \ddots & \vdots \\ K(x_t, x_1) & K(x_t, x_2) & \cdots & K(x_t, x_t) \end{pmatrix} \right)$$



Specifying Stochastic Processes

- Gaussian processes can equivalently be described via its conditional distributions:

$$f(x_{t+1})|f(x_1) = y_1, \dots, f(x_t) = y_t \\ \sim \mathcal{N}(\mu(x_{t+1}) + K_{t+1,1:t}K_{1:t,1:t}^{-1}y_{1:t}, K_{t+1,t+1} - K_{t+1,1:t}K_{1:t,1:t}^{-1}K_{1:t,t+1})$$

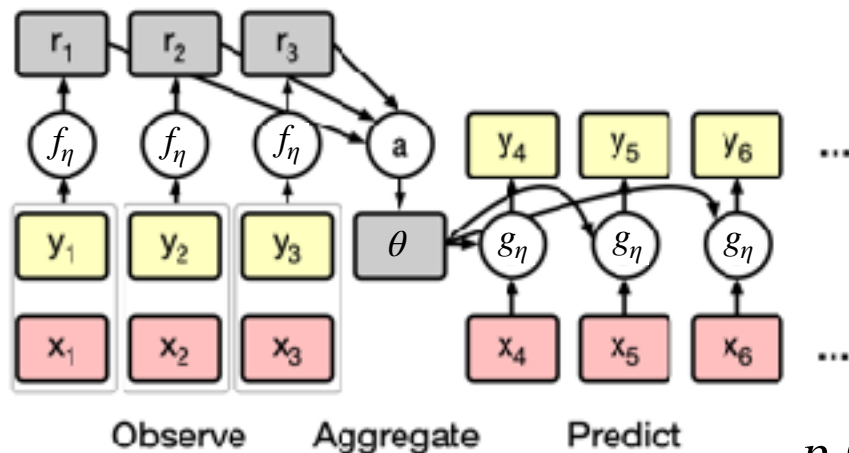
- In general, stochastic processes can also be described using a consistent family of conditional distributions:

$$p(f(x_{n+j}) = y_{n+j} | f(x_1) = y_1, \dots, f(x_n) = y_n)$$

for training dataset $\{(x_i, y_i)\}_{i=1}^n$ and test data (x_{n+j}, y_{n+j}) .



Conditional Neural Processes



- Embed input/output pairs

$$(x_i, y_i) \mapsto f_\eta(x_i, y_i)$$

- Aggregate embeddings

$$\theta = \frac{1}{n} \sum_{i=1}^n f_\eta(x_i, y_i)$$

- Predict

$$p_\eta(y_{n+j} | x_{n+j}, (x_i, y_i)_{i=1}^n) = \mathcal{N}(y_{n+j}; g_\eta(\theta, x_{n+j}))$$

- Meta-learning learns the stochastic process!



Conditional Neural Processes

Task = Function on 1D space.

Given training points, use neural processes to predict mean and std of function values at other locations.




Image Completion and Super-resolution

Task = Image = Function on 2D space.

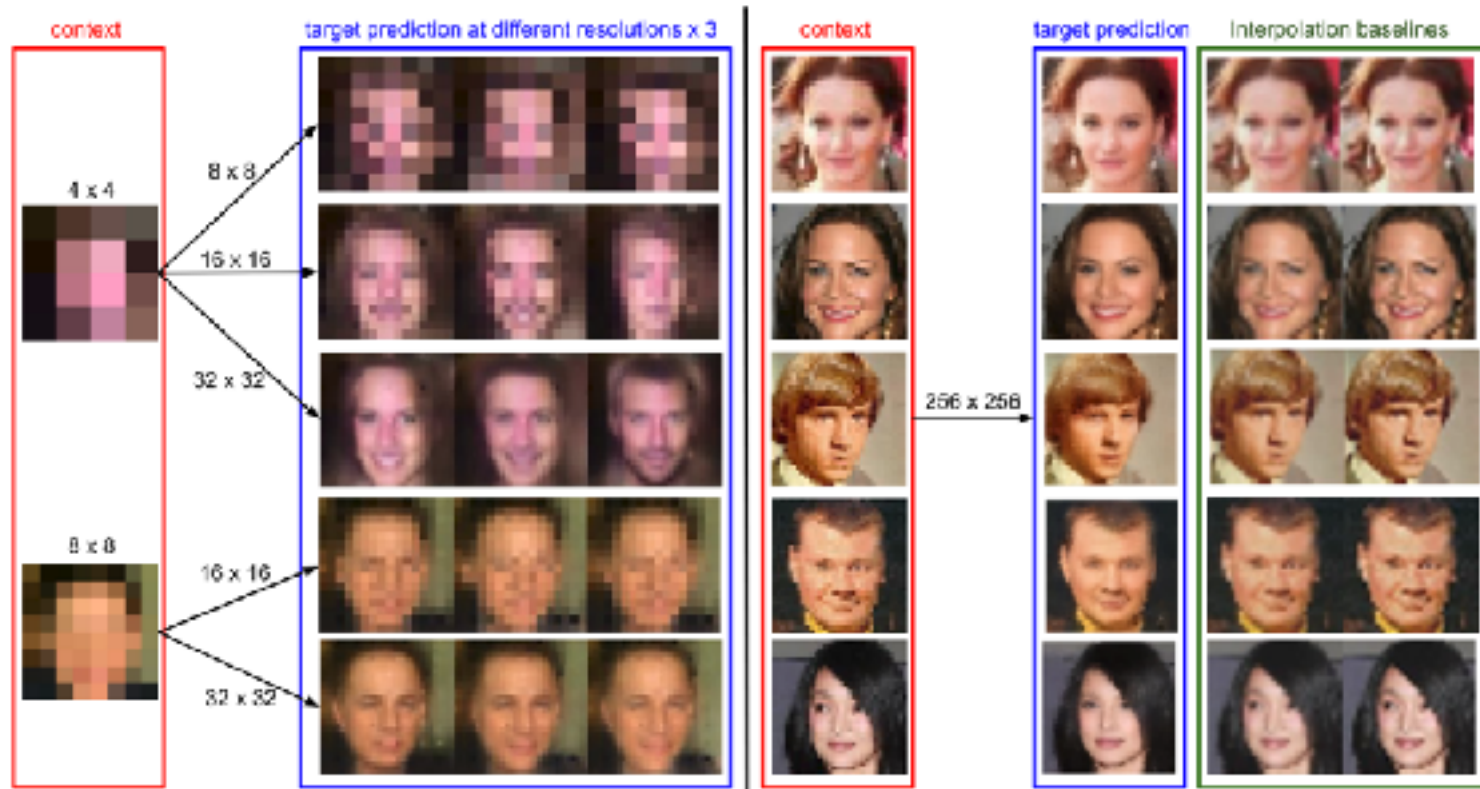
Bottom half prediction



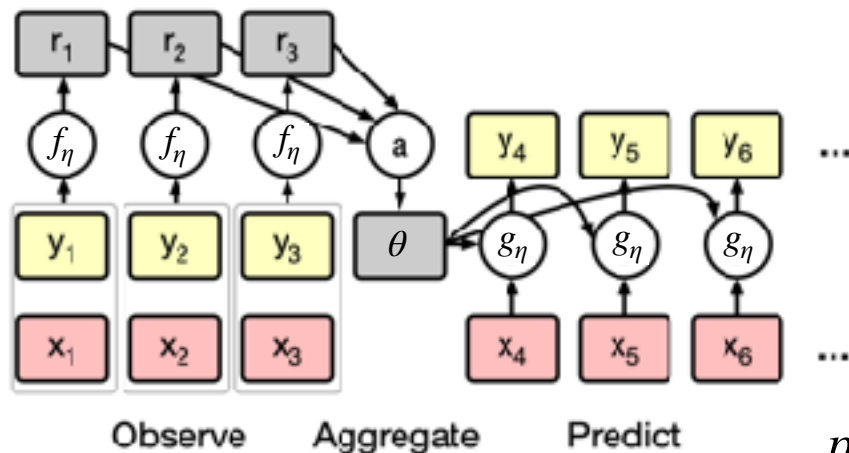
Super-resolution



Image Super-resolution



Conditional Neural Processes



- Embed input/output pairs

$$(x_i, y_i) \mapsto f_\eta(x_i, y_i)$$

- Aggregate embeddings

$$\theta = \frac{1}{n} \sum_{i=1}^n f_\eta(x_i, y_i)$$


- Predict

$$p_\eta(y_{n+j} | x_{n+j}, (x_i, y_i)_{i=1}^n) = \mathcal{N}(y_{n+j}; g_\eta(\theta, x_{n+j}))$$

- But: architecture cannot model dependence among test outputs!



Generative Process



Latent Variable Model

- Generative **model**:

$$p_{\eta}(z, y_{1:n+m} | x_{1:n+m}) = p_{\eta}(z) \prod_{i=1}^{n+m} p_{\eta}(y_i | z, x_i)$$

$$p_{\eta}(z) = \mathcal{N}(z; 0, I)$$

$$p_{\eta}(y_i | z, x_i) = \mathcal{N}(y_i; g_{\eta}(z, x_i))$$

- Variational learning objective:

$$\begin{aligned} & \log p(y_{1:n+m} | x_{1:n+m}) \\ & \geq \mathbb{E}_{q(z | x_{1:n+m}, y_{1:n+m})} \left[\sum_{i=1}^{n+m} \log p_{\eta}(y_i | z, x_i) + \log p(z) - \log q(z | x_{1:n+m}, y_{1:n+m}) \right] \\ & q(z | x_{1:t}, y_{1:t}) = \mathcal{N} \left(z; s_{\eta} \left(\frac{1}{t} \sum_{i=1}^t f_{\eta}(x_i, y_i) \right) \right) \end{aligned}$$



Alternative Learning Objective

- “our training procedure is based on a simple machine learning principle: test and train conditions must match” — [[Vinyals et al NeurIPS 2016](#)]
- Alternative objective:

$$\begin{aligned} & \log p(y_{n+:n+m} | x_{1:n}, y_{1:n}, x_{n+1:n+m}) \\ & \geq \mathbb{E}_{q(z|x_{1:n+m}, y_{1:n+m})} \left[\sum_{i=1}^{n+m} \log p_{\eta}(y_i | z, x_i) + \log p(z | x_{1:n}, y_{1:n}) - \log q(z | x_{1:n+m}, y_{1:n+m}) \right] \\ & \approx \mathbb{E}_{q(z|x_{1:n+m}, y_{1:n+m})} \left[\sum_{i=1}^{n+m} \log p_{\eta}(y_i | z, x_i) + \log q(z | x_{1:n}, y_{1:n}) - \log q(z | x_{1:n+m}, y_{1:n+m}) \right] \\ & q(z | x_{1:t}, y_{1:t}) = \mathcal{N} \left(z; s_{\eta} \left(\frac{1}{t} \sum_{i=1}^t f_{\eta}(x_i, y_i) \right) \right) \end{aligned}$$




Probabilistic Perspective on Meta-Learning

- Meta-learning as learning a stochastic process:
 - Meta-learning a prior over functions -> meta-learning inductive biases from meta-training set!
 - Base learner can be thought of as **amortized learning**.
- Uncertainties are important in meta-learning applications:
 - Active learning
 - Bayesian optimisation
 - Reinforcement learning



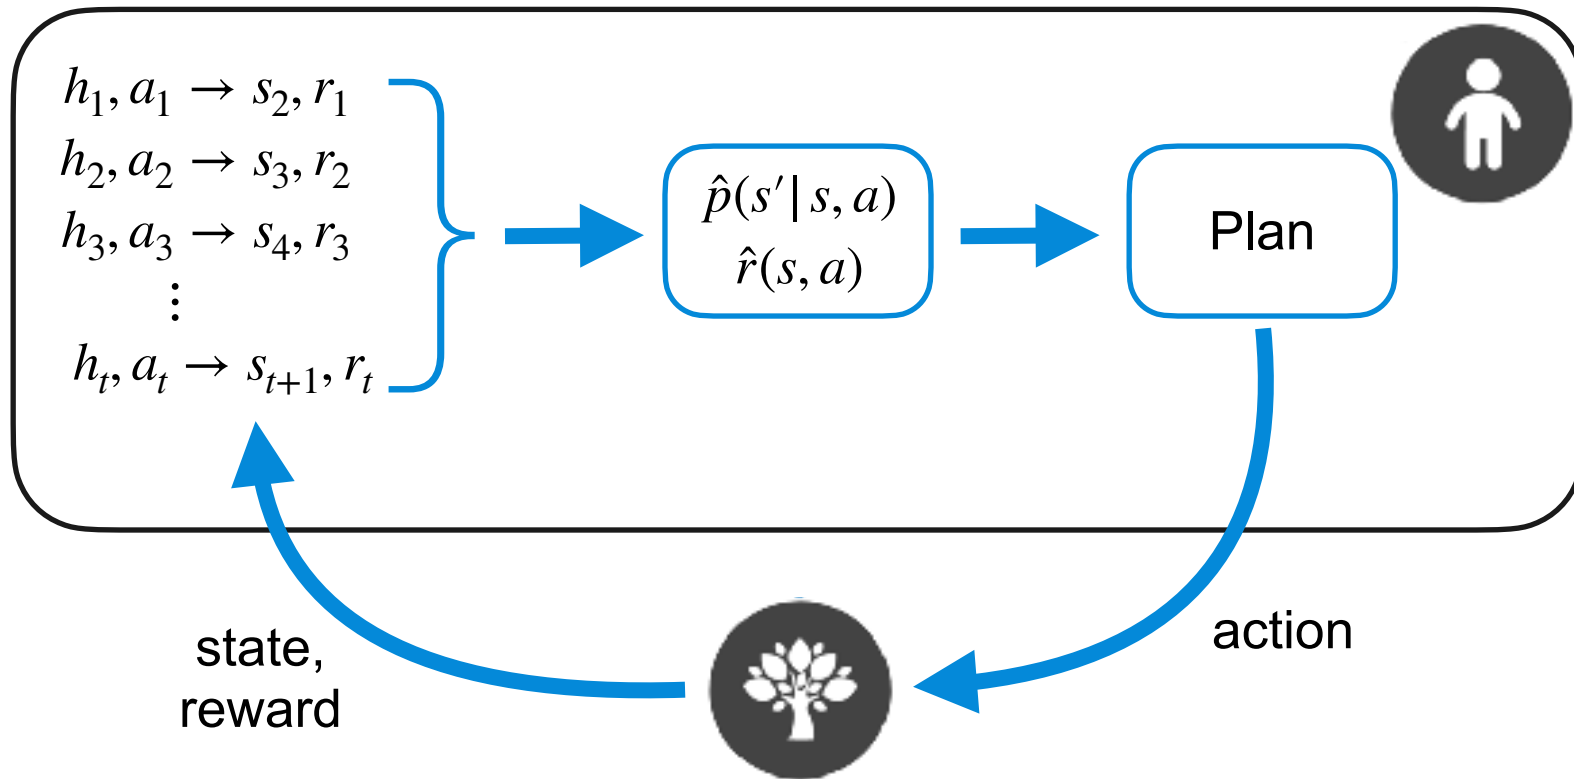
Uncertainty in Meta-Learning




- Small training sets in meta-learning \rightarrow uncertainty in task inference!
- Important for:
 - active learning,
 - Bayesian optimisation,
 - reinforcement learning



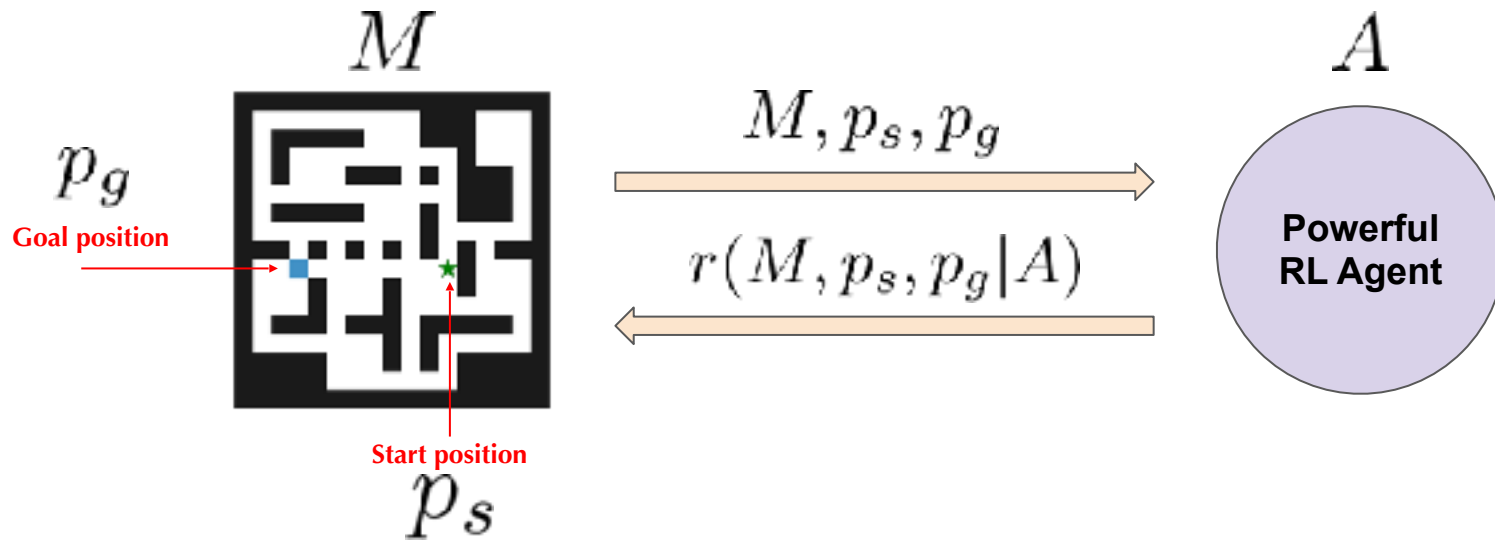
Efficient Model-based Reinforcement Learning




Cart Pole



Adversarial Testing of RL Agents



Adversarial Testing of RL Agents




Bayesian
Optimization

$$\min_{M, p_s, p_g} r(M, p_s, p_g | A)$$


$(M, p_s, p_g, A) \sim p(\mathcal{T})$ - training & holdout samples (agents, mazes, positions)



Adversarial Testing of RL Agents



Bayesian optimisation iterations

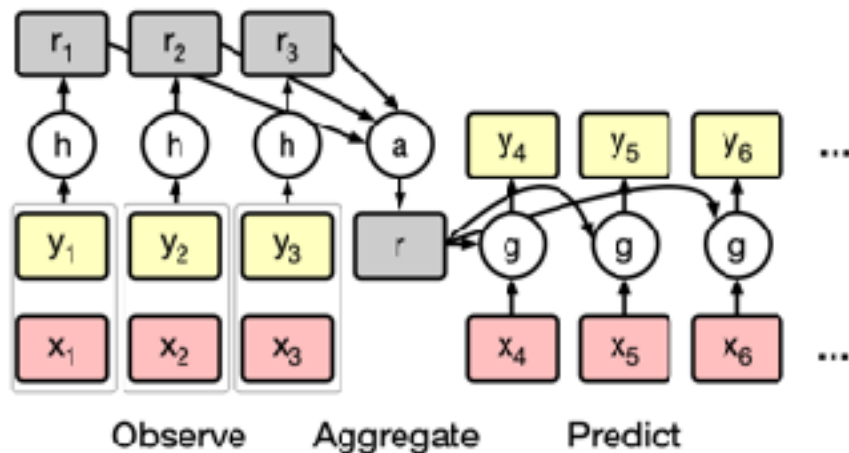


Meta-Learning: an idiosyncratic tutorial

- Optimisation perspective on meta-learning
 - Optimisation-based meta-learning
 - Black-box meta-learning
- Probabilistic perspective on meta-learning
 - Stochastic processes
 - Neural processes
 - Uncertainty in meta-learning
- **Probabilistic symmetries and neural architectures**
- Note: no meta reinforcement learning (meta-RL)



Permutation-Invariance in Neural Processes



- Embed input/output pairs

$$(x_i, y_i) \mapsto f_\eta(x_i, y_i)$$

- Aggregate embeddings

$$\theta = \frac{1}{n} \sum_{i=1}^n f_\eta(x_i, y_i)$$

- Predict

$$p_\eta(y_{n+j} | x_{n+j}, (x_i, y_i)_{i=1}^n) = \mathcal{N}(y_{n+j}; g_\eta(\theta, x_{n+j}))$$



Characterising Permutation-Invariant Functions

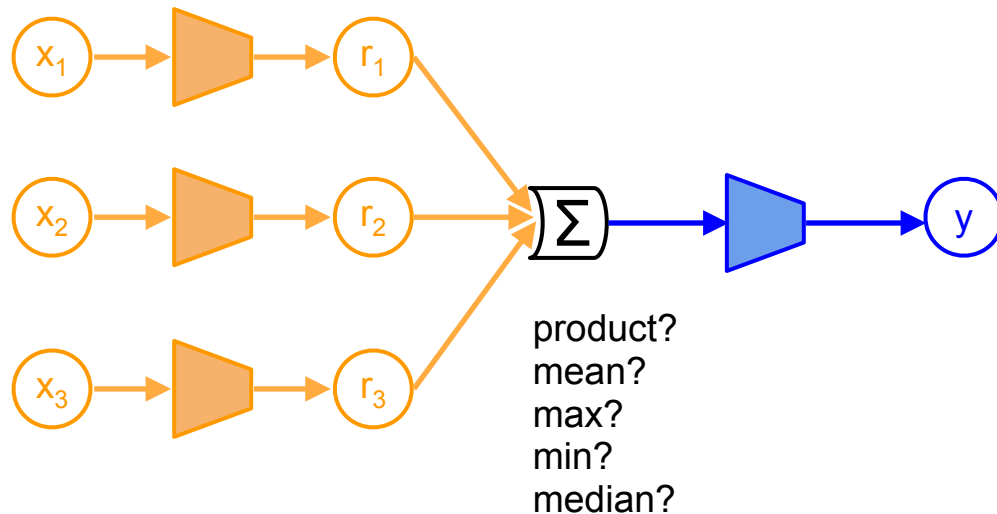
- Function $h : \mathcal{X}^n \rightarrow \mathcal{Y}$ is permutation-invariant,

$$h(\pi \cdot (x_1, \dots, x_n)) = h(x_{\pi(1)}, \dots, x_{\pi(n)}) = h(x_1, \dots, x_n)$$

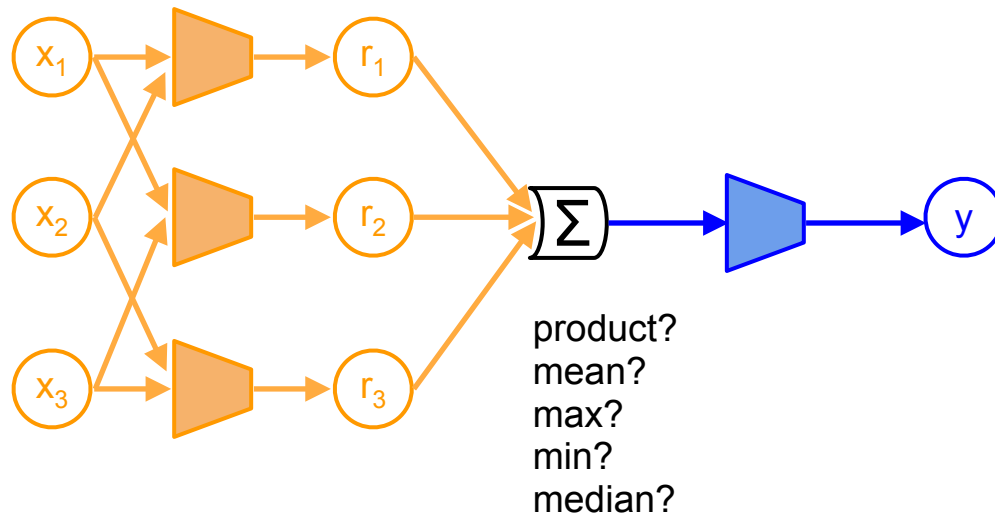
- Can we characterise the class of permutation-invariant functions?
- If we use neural networks to parameterise permutation-invariant functions, how should we choose the architecture?
- Given an architecture choice, can the neural network approximate well any arbitrary permutation-invariant function?



Characterising Permutation-Invariant Functions



Characterising Permutation-Invariant Functions



Functional Symmetry Properties

- Function $h : \mathcal{X}^n \rightarrow \mathcal{Y}^n$ is **permutation-equivariant**,

$$h(x_1, \dots, x_n) = (y_1, \dots, y_n)$$

$$h(\pi \cdot (x_1, \dots, x_n)) = \pi \cdot (y_1, \dots, y_n) = \pi \cdot h(x_1, \dots, x_n)$$

- **Group** G acting on input space \mathcal{X} and output space \mathcal{Y} .

- **G -invariant:**

$$h(g \cdot x) = h(x)$$

- **G -equivariant:**

$$h(g \cdot x) = g \cdot h(x)$$



Probabilistic Symmetries

- A distribution P for a random sequence $\mathbf{X}_n = (X_1, \dots, X_n)$ is **exchangeable** if

$$P(X_1, \dots, X_n) = P(\pi \cdot (X_1, \dots, X_n))$$

$$P(X_1 \in B_1, \dots, X_n \in B_n) = P(X_{\pi(1)} \in B_1, \dots, X_{\pi(n)} \in B_n))$$

- Exchangeability is permutation-invariance of P .
- $\mathbf{X}_{\mathbb{N}}$ is infinitely exchangeable if all length n prefixes are exchangeable.
- **de Finetti's Theorem:**

$\mathbf{X}_{\mathbb{N}}$ is infinitely exchangeable $\Leftrightarrow X_i | Q \sim_{iid} Q$ for some random Q .



Probabilistic Symmetries for Conditional Distributions

- A conditional distribution $P(Y|X)$ is a stochastic relaxation for a function $Y = h(X)$.

- $P(Y|X)$ is G -invariant if:

$$P(Y|X) = P(Y|g \cdot X)$$

$$P(Y \in B|X \in A) = P(Y \in B|g \cdot X \in A)$$

- $P(Y|X)$ is G -equivariant if:

$$P(Y|X) = P(g \cdot Y|g \cdot X)$$

- Can we characterise the class of permutation-invariant conditional distributions?



Empirical Measure

- de Finetti's Theorem may fail for finitely exchangeable sequences.
- The **empirical measure** of \mathbf{X}_n is

$$\mathbb{M}_{\mathbf{X}_n}(\cdot) = \sum_{i=1}^n \delta_{X_i}(\cdot)$$

- The empirical measure is a **sufficient statistic**: P is exchangeable iff

$$P(\mathbf{X}_n \in \cdot \mid \mathbb{M}_{\mathbf{X}_n} = m) = \mathbb{U}_m(\cdot)$$

where \mathbb{U}_m is the uniform distribution over all sequences (x_1, \dots, x_n) with empirical measure m .



Noise Outsourcing

- If X and Y are random variables in “nice” (e.g. Borel) spaces \mathcal{X} and \mathcal{Y} , then there are a random variable $\eta \sim U[0,1]$ with $\eta \perp\!\!\!\perp X$ and a function $h : [0,1] \times \mathcal{X} \mapsto \mathcal{Y}$ such that

$$(X, Y) =_{a.s.} (X, h(\eta, X))$$

- Furthermore, if there is an **adequate statistic** $S(X)$ with $X \perp\!\!\!\perp Y | S(X)$, then

$$(X, Y) =_{a.s.} (X, h(\eta, S(X)))$$



Probabilistic Permutation-Invariance

- Now suppose we have random variables \mathbf{X}_n and Y .
 - Y is conditionally permutation-invariant given \mathbf{X}_n .
 - \mathbf{X}_n is marginally permutation-invariant (exchangeable).
- The empirical measure is a sufficient statistic for \mathbf{X}_n .
- It is also an adequate statistic for Y given \mathbf{X}_n :

$$P(Y | \mathbf{X}_n = \mathbf{x}_n) = P(Y | \mathbb{M}_{\mathbf{X}_n} = \mathbb{M}_{\mathbf{x}_n})$$


We have the conditional independence $\mathbf{X}_n \perp\!\!\!\perp Y | \mathbb{M}_{\mathbf{X}_n}$.

- Noise outsourcing...


$$(\mathbf{X}_n, Y) =_{a.s.} (\mathbf{X}_n, h(\eta, \mathbb{M}_{\mathbf{X}_n}))$$




Probabilistic Permutation-Invariance



Probabilistic Permutation-Invariance



Functional Permutation-Invariance



Probabilistic Permutation-Equivariance

- Now suppose we have random sequences \mathbf{X}_n and \mathbf{Y}_n .
 - \mathbf{Y}_n is conditionally permutation-equivariant given \mathbf{X}_n .
 - \mathbf{X}_n is marginally permutation-invariant (exchangeable).
- Also suppose that $Y_i \perp\!\!\!\perp \mathbf{Y}_n \setminus Y_i \mid \mathbf{X}_n$ for each i .


- Then:

$$(\mathbf{X}_n, (Y_1, \dots, Y_n)) =_{a.s.} (\mathbf{X}_n, (h(\eta_1, X_1, \mathbb{M}_{\mathbf{X}_n}), \dots, h(\eta_n, X_n, \mathbb{M}_{\mathbf{X}_n})))$$



for outsourced noise (η_i) that are identical, mutually independent and independent of \mathbf{X}_n .



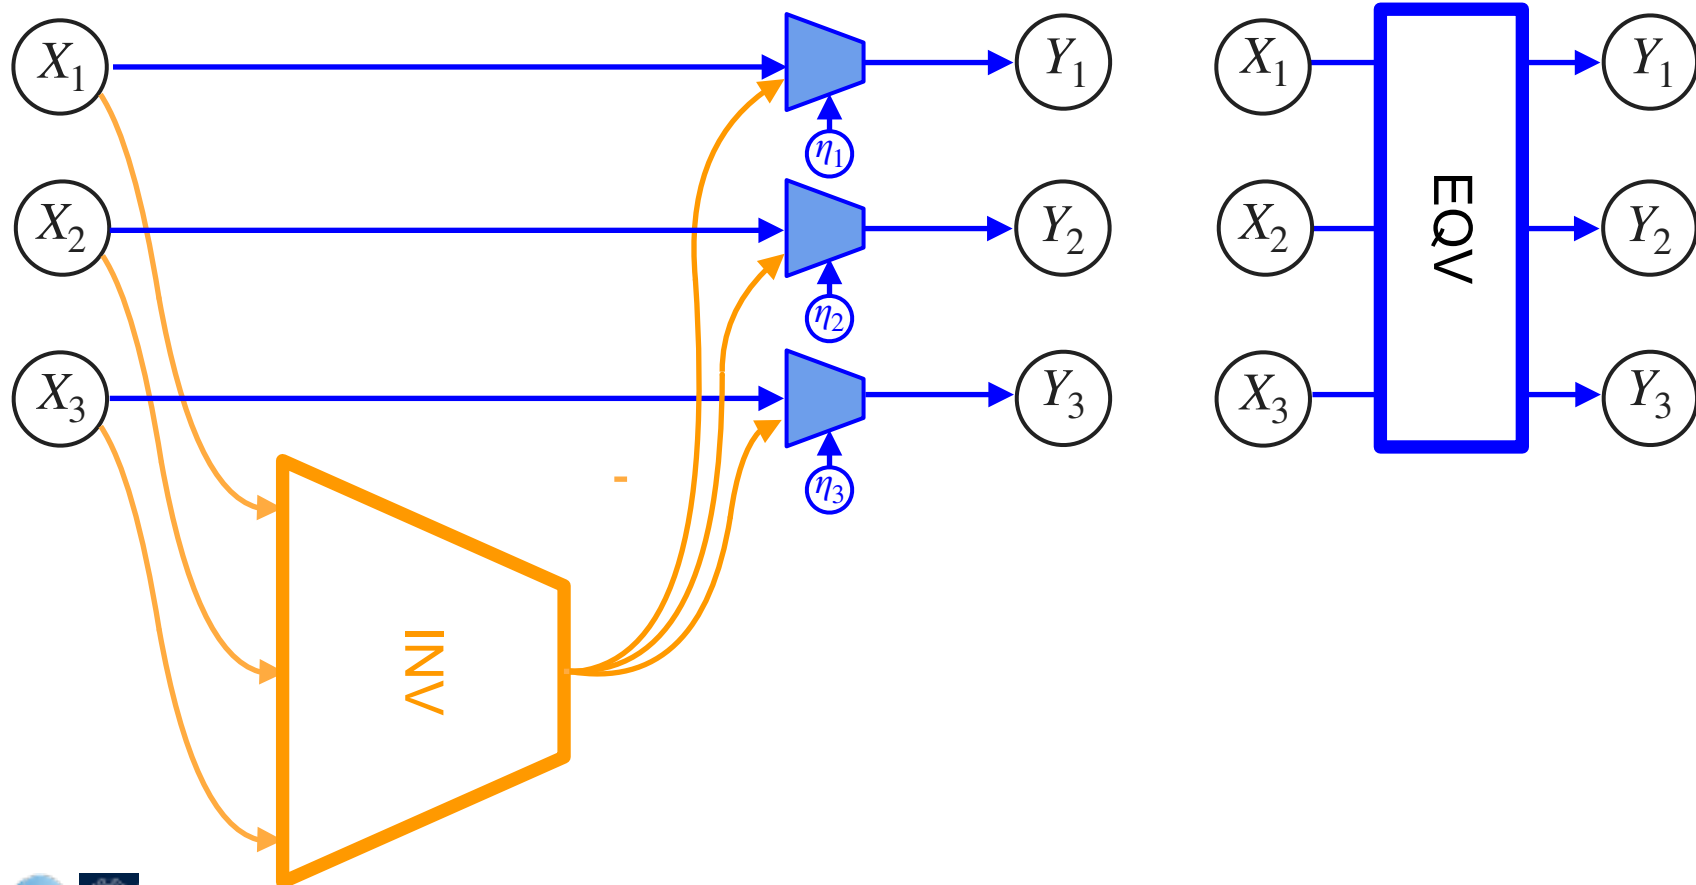
Probabilistic Permutation-Equivariance




Composing Invariant and Equivariant Modules




Composing Invariant and Equivariant Modules



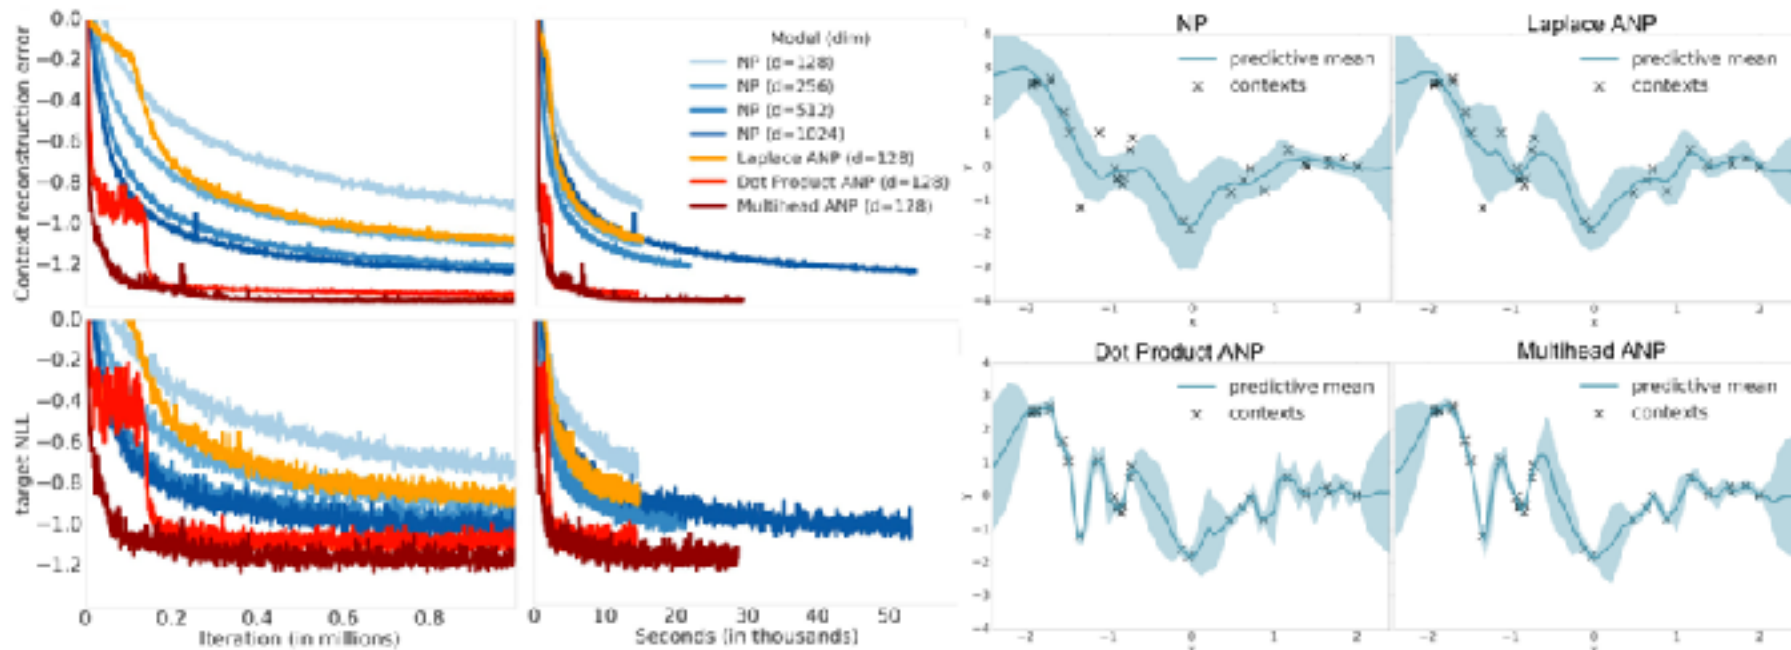
Composing Invariant and Equivariant Modules



Composing Invariant and Equivariant Modules



Attentive Neural Processes



Maximal Invariant and Maximal Equivariant

- Let G be a compact group.
- A maximal invariant is a statistic $M : \mathcal{X} \mapsto \mathcal{S}$ such that

$$M(g \cdot x) = M(x) \forall g \in G, x \in \mathcal{X}$$

$$M(x_1) = M(x_2) \Rightarrow \exists g \in G : x_1 = g \cdot x_2$$

- A maximal equivariant $\tau : \mathcal{X} \mapsto G$ satisfies

$$\tau(g \cdot x) = g \cdot \tau(x)$$



Probabilistic and Functional Symmetries

- Let G be a compact group and X be marginally G -invariant.

- Let M be a maximal invariant, then

Y is conditionally G -invariant given $X \Leftrightarrow (X, Y) =_{a.s.} (X, h(\eta, M(X)))$

for outsourced noise η independent of X and a function h .

- If a maximal equivariant τ exists and $G_X \subset G_Y$ a.s., then

Y is conditionally G -equivariant given $X \Leftrightarrow (X, Y) =_{a.s.} (X, h(\eta, X))$

for outsourced noise η independent of X and a function h that is G -equivariant in its second argument.



Probabilistic Symmetries

- Tools from probabilistic symmetry, sufficiency and adequacy allowed us to answer questions about neural architectures under symmetry.
 - Framework extends to graph and array structured data with node exchangeability.
 - Continuous groups?
 - How to relax assumptions of conditional independence of outputs?



Meta-Learning: an idiosyncratic tutorial

- Optimisation perspective on meta-learning
 - Optimisation-based meta-learning
 - Black-box meta-learning
- Probabilistic perspective on meta-learning
 - Stochastic processes
 - Neural processes
 - Uncertainty in meta-learning
- Probabilistic symmetries and neural architectures
- Note: no meta reinforcement learning (meta-RL)



Some Recent Developments

- Recent Developments:
 - [[Learning to Continually Learn](#)]
 - [[MetaFun: Meta-Learning with Iterative Functional Updates](#)]
 - [[Meta-Gradient Reinforcement Learning](#)]
 - [[Modular Meta-Learning with Shrinkage](#)]
 - [[Rapid Learning or Feature Reuse? Towards Understanding the Effectiveness of MAML](#)]
 - [[Meta-Learning with Warped Gradient Descent](#)]
- Meta-learning for continual learning
 - [[Learning to Continually Learn](#)]
 - [[Meta-Learning Representations for Continual Learning](#)]
 - [[Task Agnostic Continual Learning via Meta Learning](#)]
- Meta-learning theory
 - [[Theoretical Models of Learning to Learn](#)]
 - [[A Model of Inductive Bias Learning](#)]
 - [[Provably Meta-Learning of Linear Representations](#)]



Thank You!

- Collaborators, colleagues, mentors



- You!



<http://csml.stats.ox.ac.uk/people/teh>



Yee Whye Teh